

PSTAT 174 Project Data Tidy

Mu Niu

2023-05-09

Data Preprocessing

(1) We first read in the csv file to get the raw data:

```
raw <- read.csv('GlobalLandTemperaturesByMajorCity.csv')
head(raw) #Look at the raw data
```

```
##           dt AverageTemperature AverageTemperatureUncertainty   City
## 1 1849-01-01           26.704                1.435 Abidjan
## 2 1849-02-01           27.434                1.362 Abidjan
## 3 1849-03-01           28.101                1.612 Abidjan
## 4 1849-04-01           26.140                1.387 Abidjan
## 5 1849-05-01           25.427                1.200 Abidjan
## 6 1849-06-01           24.844                1.402 Abidjan
##           Country Latitude Longitude
## 1 Côte D'Ivoire    5.63N    3.23W
## 2 Côte D'Ivoire    5.63N    3.23W
## 3 Côte D'Ivoire    5.63N    3.23W
## 4 Côte D'Ivoire    5.63N    3.23W
## 5 Côte D'Ivoire    5.63N    3.23W
## 6 Côte D'Ivoire    5.63N    3.23W
```

(2) My goal is to forecast the temperature of Peking, the capital of China, so I have to filter the data set:

```
pek <- raw[raw$City == "Peking",]
head(pek)
```

```
##           dt AverageTemperature AverageTemperatureUncertainty   City
## 176248 1820-08-01           22.822                2.218 Peking
## 176249 1820-09-01           19.384                1.706 Peking
## 176250 1820-10-01           11.029                1.817 Peking
## 176251 1820-11-01            2.429                1.961 Peking
## 176252 1820-12-01           -4.767                2.237 Peking
## 176253 1821-01-01           -3.565                2.213 Peking
```

```
##          Country Latitude Longitude
## 176248    China   39.38N   116.53E
## 176249    China   39.38N   116.53E
## 176250    China   39.38N   116.53E
## 176251    China   39.38N   116.53E
## 176252    China   39.38N   116.53E
## 176253    China   39.38N   116.53E
```

- Since we want to work on univariate time series, we only want to keep the variable that we want to forecast, which is the Average Temperature:

```
pek.temp <- pek[,1:2]
head(pek.temp)
```

```
##          dt AverageTemperature
## 176248 1820-08-01          22.822
## 176249 1820-09-01          19.384
## 176250 1820-10-01          11.029
## 176251 1820-11-01           2.429
## 176252 1820-12-01          -4.767
## 176253 1821-01-01          -3.565
```

(3) Now, we want to check whether there is missing values:

```
sum(is.na(pek.temp))
```

```
## [1] 14
```

Knowing there are 14 missing values, we want to know how the missing values are distributed(index of NaN):

```
pek.temp[is.na(pek.temp$AverageTemperature) == TRUE,]
```

```
##          dt AverageTemperature
## 176394 1832-10-01             NA
## 176457 1838-01-01             NA
## 176458 1838-02-01             NA
## 176459 1838-03-01             NA
## 176460 1838-04-01             NA
## 176461 1838-05-01             NA
## 176462 1838-06-01             NA
## 176463 1838-07-01             NA
## 176464 1838-08-01             NA
## 176465 1838-09-01             NA
## 176466 1838-10-01             NA
## 176467 1838-11-01             NA
## 176468 1838-12-01             NA
## 178565 2013-09-01             NA
```

We notice that we don't have observation for October 1832, all year of 1838, and September 2013.

(4) Fill the Missing Values:

- For October 1832: Take the mean value of the neighbors average temperature
- For September 2013: This is the last observation in our data set, so we can't take the mean value of the neighbors average temperature. What we want to do is take the mean value of the two previous months' average temperature, since July, August, and September are in the same quarter of the year.
- For all 12 months of 1838: Even the data set is monthly, I want to make it yearly data by taking the average temperature of 12 months. Hence, we can first fill all 12 months in 1838 by 0. After we get the yearly data, we can change it to the mean value of the 2 neighbor year average temperature.

We first fill the value of October 1832, September 2013, and assign 0 to 1838 all year:

```
pek.temp["176394", "AverageTemperature"] <- (pek.temp["176395", "AverageTemperature"]+pek.temp["176393", "AverageTemperature"])/2
pek.temp["178565", "AverageTemperature"] <- (pek.temp["178563", "AverageTemperature"]+pek.temp["178564", "AverageTemperature"])/2
for (index in as.character(seq(176457,176468))) {
  pek.temp[index, "AverageTemperature"] <- 0
}

# Check Missing Values
sum(is.na(pek.temp))
```

```
## [1] 0
```

(5) Convert to yearly data:

Looking at the data set, I noticed that both 1820 and 2013 don't have observations for all 12 months. Considering the lack of particular months might lead to mistake, I decided to drop these years.

```
pek.temp <- pek.temp[as.character(seq(176253,178556)),] # Monthly temperature data from 1821 to 2012
```

Then, we convert the monthly data to yearly data by taking average.(yearly data is more smooth than monthly data, and ideal to work with):

```
year.temp <- c()
for (jan in seq(1,2304,12)){
  dec <- jan+11
  year.temp <- c(year.temp,mean(pek.temp[jan:dec,2]))
}
```

Now, we can construct a new data frame by using the yearly data, and fill the value of 1838 by using the mean of its neighbors:

```
Year <- c(1821:2012)
data <- data.frame(Year = Year,
                   AverageTemperature = year.temp)
data[data$Year == 1838, "AverageTemperature"] <- (data[data$Year == 1837, "AverageTemperature"] +
                                                    data[data$Year == 1839, "AverageTemperature"])/2
head(data)
```

```
##   Year AverageTemperature
## 1 1821          11.50625
## 2 1822          11.50225
## 3 1823          10.96133
## 4 1824          11.96475
## 5 1825          11.73483
## 6 1826          11.65225
```

We need split the data into two parts: we use the first part to model the data, and the the second part to test our model

```
model.data <- data[Year <= 2010, ]
test.data <- data[Year > 2010, ]
```

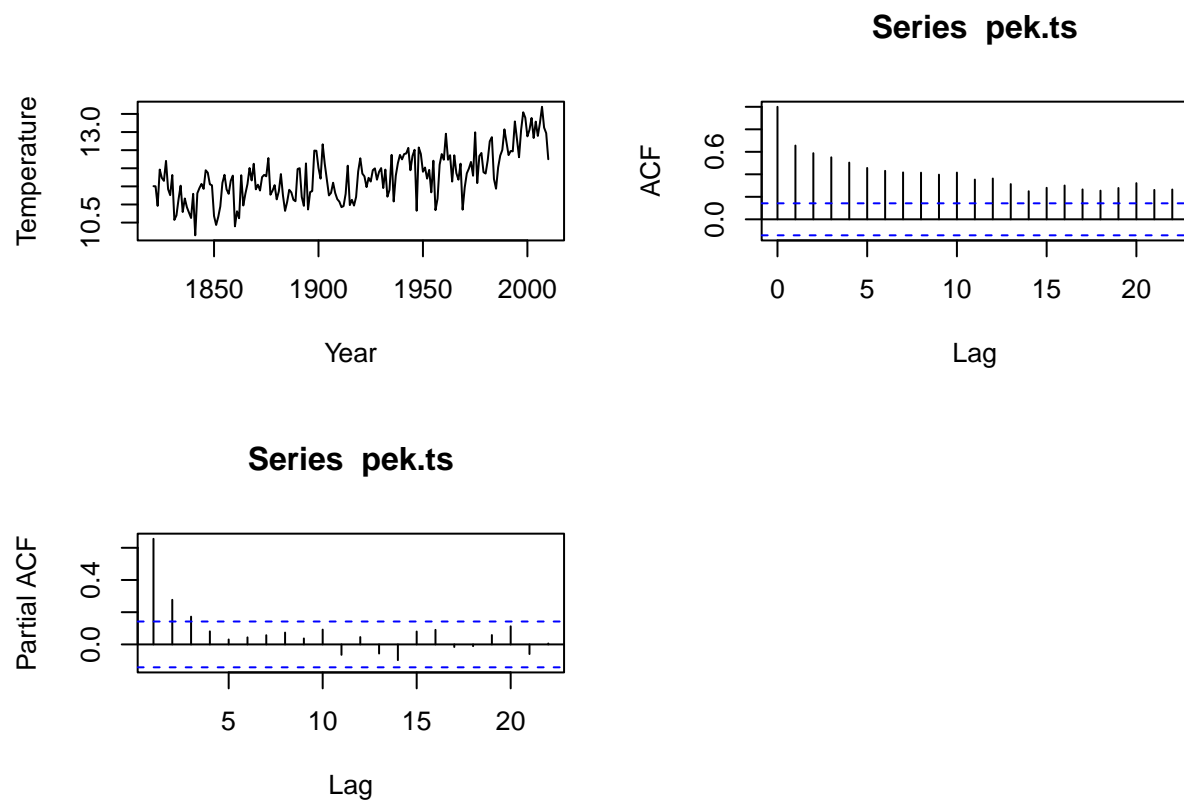
Now, we finished the process of data tidying, and can use this data for time series study.

```
# Export the data set
write.csv(data, "Peking.Temp.1821-2012.csv")
write.csv(model.data, "train.1821-2011.csv")
write.csv(test.data, "test.2012.csv")
```

Transformation and Differencing

Original Peking temperature data set from 1821-2011:

```
op <- par(mfrow=c(2,2))
pek.ts <- ts(model.data[,2], start = 1821, frequency = 1)
ts.plot(pek.ts, gpars=list(xlab="Year", ylab="Temperature"))
acf(pek.ts)
pacf(pek.ts)
par(op)
```




From the plots above, we can see that there is no seasonality but an upward trend. The variance is relatively stable, but we need to check the value of λ and the histograms to see if transformation is necessary.

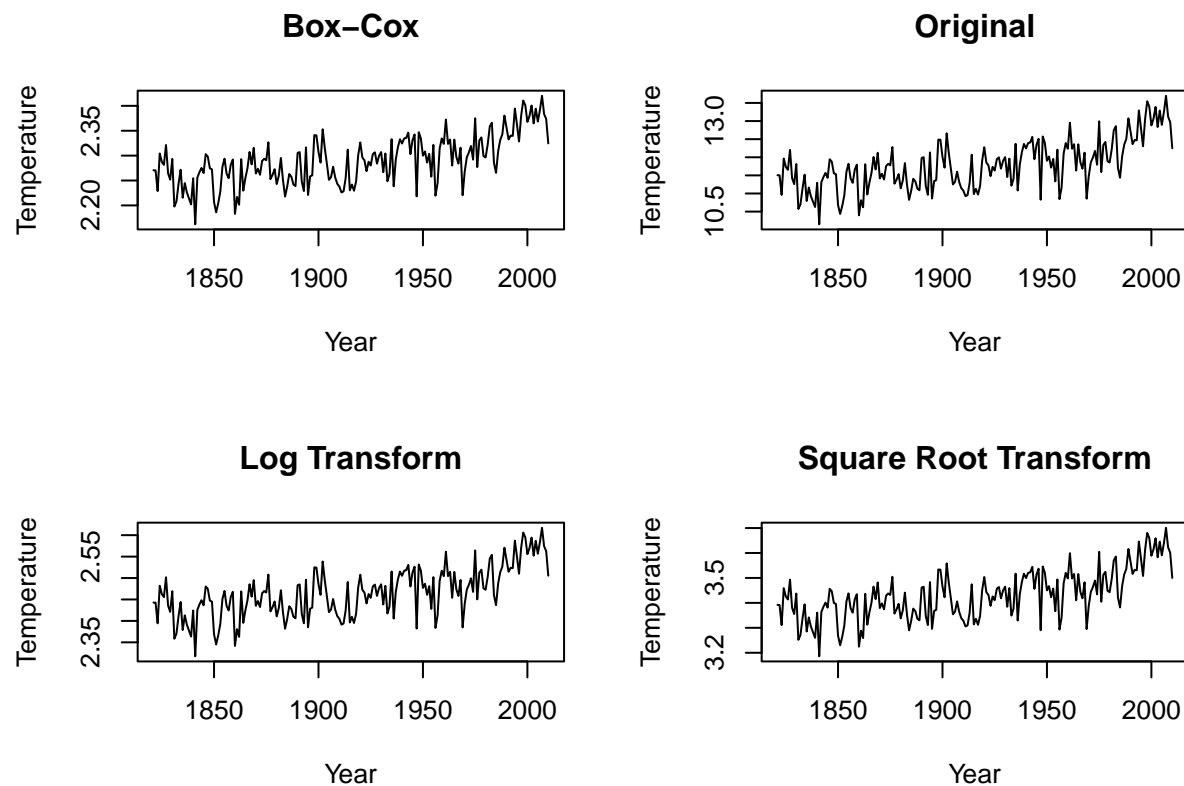
Transformation

Now, let's see if we need any transformations:

```
library(MASS)
t = 1:length(pek.ts)
fit = lm(pek.ts ~ t)
bcTransform = boxcox(pek.ts ~ t, plotit = T)
```

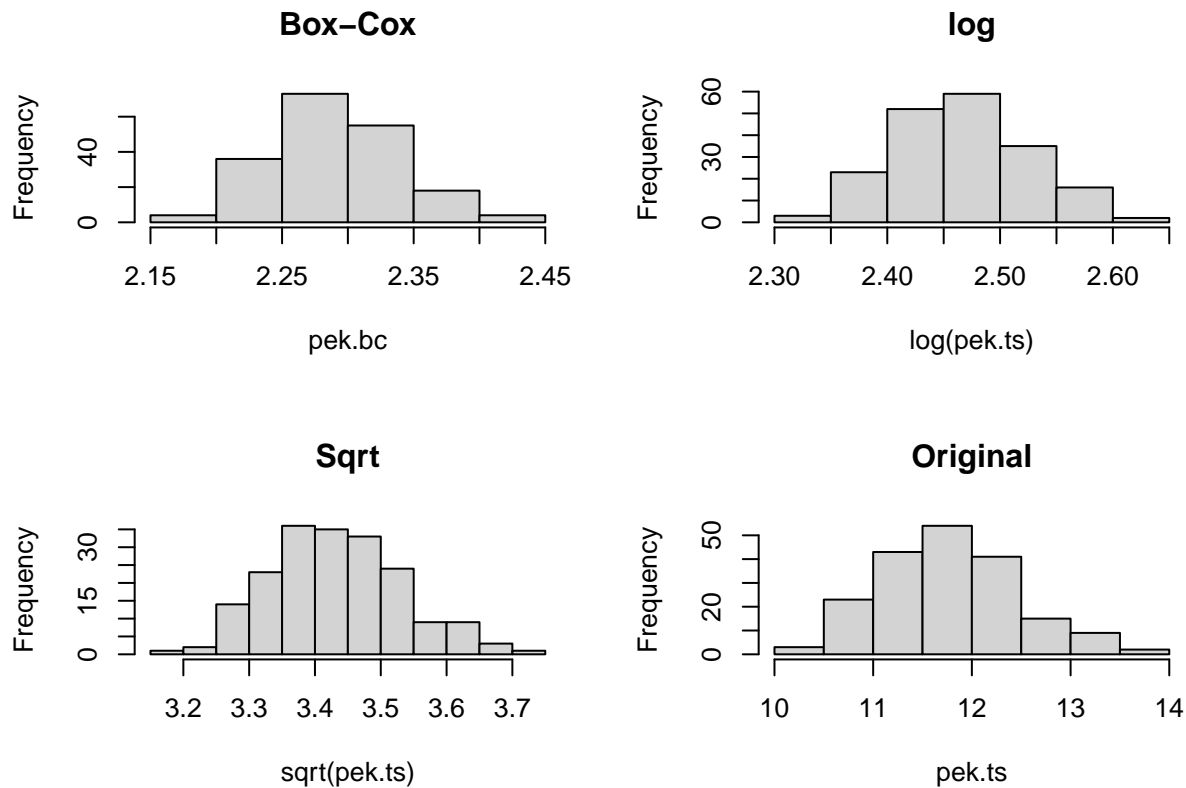


```
lambda = bcTransform$x[which(bcTransform$y == max(bcTransform$y))]
pek.bc = (1/lambda)*(pek.ts^lambda-1)
op <- par(mfrow=c(2,2))
ts.plot(pek.bc,main = "Box-Cox", xlab='Year',ylab='Temperature')
ts.plot(pek.ts,main = "Original", xlab='Year',ylab='Temperature')
ts.plot(log(pek.ts), main = "Log Transform", xlab='Year',ylab='Temperature')
ts.plot(sqrt(pek.ts), main = "Square Root Transform", xlab='Year',ylab='Temperature')
```



```
par(op)
```

```
op <- par(mfrow=c(2,2))
hist(pek.bc, main = 'Box-Cox')
hist(log(pek.ts),main='log')
hist(sqrt(pek.ts),main= 'Sqrt')
hist(pek.ts, main = 'Original')
```



```
par(op)
```

The Box-Cox plot shows that 1 is within the 95% Confidence Interval for the value of λ , also the histogram, which tells us the distribution of the data before transformation, appears to be normally distributed and has a bell shape. Hence, it's safe to conclude that our data doesn't need to be transformed.

Differencing

Since we noticed the upward trend, let's first difference the data once at lag 1 and compare its variance with the variance of data before differencing:

```
var(pek.ts) # Before Diff
```

```
## [1] 0.4770113
```

```
dt.pek.ts <- diff(pek.ts, 1)
(var(dt.pek.ts)) # Diff once at lag 1
```

```
## [1] 0.328638
```

Then, let's difference again at lag 1 and compare their variance:


```
dt2.pek.ts <- diff(dt.pek.ts, 1)
(var(dt2.pek.ts)) # Diff twice at lag 1
```


```
## [1] 0.9280505
```

We can see the variance increased after we differentiate twice, which means over differencing. Hence, we want to use the diff once data.

Model Identification


We first want to guess the model based on the ACF and PACF:

```
plot(dt.pek.ts)
```




```
acf(dt.pek.ts, lag.max = 40)
```

Series dt.pek.ts



```
pacf(dt.pek.ts, lag.max = 40)
```

Series dt.pek.ts



Now, we can fit different ARMA models using maximum likelihood estimation and compare the model fits using AICC:

- p: 0,1,2,3
- d: 1
- q: 0,1

```
library(qpcR)
```

```
## Loading required package: minpack.lm
```

```
## Loading required package: rgl
```

```
## Loading required package: robustbase
```

```
## Loading required package: Matrix
```

```
# Construct Matrix
```

```
aiccs <- matrix(NA, nr = 4, nc = 2)  
dimnames(aiccs) = list(p=0:3, q=0:1)
```

```
# Use for loop to calculate the AICc matrix
```

```
for(p in 0:3) {  
  for(q in 0:1) {
```

```

    aiccs[p+1,q+1] = AICc(arima(dt.pek.ts, order = c(p,0,q), method="ML"))
  } }
aiccs

```

```

##      q
## p      0      1
## 0 329.0586 270.5249
## 1 296.4201 265.7672
## 2 284.8366 265.8102
## 3 281.9829 266.3851

```

Another way to identify the model is to use the automatic arima fit function by Rob Hyndman:

```
library(forecast)
```

```

## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo

```

```
auto.arima(pek.ts)
```

```

## Series: pek.ts
## ARIMA(1,1,1) with drift
##
## Coefficients:
##          ar1      ma1    drift
##          0.2847 -0.9210  0.0088
## s.e.    0.0942   0.0512  0.0041
##
## sigma^2 = 0.2309: log likelihood = -128.82
## AIC=265.64   AICc=265.85   BIC=278.6

```

We can choose the models that has low AICc from the table: *ARIMA*(1,1,1) and *ARIMA*(2,1,1)

Coefficients Estimation and Diagnostic Checking

- Coefficients Estimation

Now, we can fit the data in the chosen models, which is *ARIMA*(1,1,1) and *ARIMA*(2,1,1):

- *ARIMA*(1,1,1)

```
(arima111 <- arima(pek.ts, order=c(1,1,1), method="ML"))
```

```
##
## Call:
## arima(x = pek.ts, order = c(1, 1, 1), method = "ML")
##
## Coefficients:
##          ar1      ma1
##      0.2373 -0.8656
## s.e.  0.0963  0.0547
##
## sigma^2 estimated as 0.2308:  log likelihood = -130.12,  aic = 266.24
```

- *ARIMA*(2,1,1)

```
(arima211 <- arima(pek.ts, order=c(2,1,1), method="ML"))
```

```
##
## Call:
## arima(x = pek.ts, order = c(2, 1, 1), method = "ML")
##
## Coefficients:
##          ar1      ar2      ma1
##      0.2519  0.0870 -0.8949
## s.e.  0.0890  0.0839  0.0479
##
## sigma^2 estimated as 0.2295:  log likelihood = -129.59,  aic = 267.18
```

Since the 95% Confidence Interval of ϕ_2 contains 0, we can set it to zero. But doing so will give us *ARIMA*(1,1,1) model, which is exactly the same with the previous one. So we remove *ARIMA*(2,1,1) and consider *MA*(1) with AICc = 270.5249 and *AR*(3) with AICc = 281.9829, because they have low AICc in comparison with other models:

```
(ma1 <- arima(pek.ts, order=c(0,1,1), method="ML"))
```

```
##
## Call:
## arima(x = pek.ts, order = c(0, 1, 1), method = "ML")
##
## Coefficients:
##          ma1
##      -0.7186
## s.e.  0.0742
##
## sigma^2 estimated as 0.2371:  log likelihood = -132.53,  aic = 269.07
```

```
(ar3 <- arima(pek.ts, order=c(3,1,0), method="ML"))
```

```
##
## Call:
## arima(x = pek.ts, order = c(3, 1, 0), method = "ML")
##
## Coefficients:
```

```
##          ar1          ar2          ar3
##      -0.5644  -0.3494  -0.1620
## s.e.   0.0724   0.0793   0.0728
##
## sigma^2 estimated as 0.2463:  log likelihood = -135.96,  aic = 279.91
```

Since $MA(1)$ has fewer parameters and lower AICc value ($270.5249 < 281.9829$), it's better to only take the $MA(1)$ model and $ARIMA(1,1,1)$ to diagnostic checking.

- Diagnostic Checking

Let's first write out the two models:

Let X_t denotes our data,

(A) $ARIMA(1,1,1)$:

$$(1 - 0.2373_{0.0963}B)\nabla_1 X_t = (1 - 0.8656_{0.0547}B)Z_t \text{ with } \hat{\sigma}_Z^2 = 0.2308$$

(B) $MA(1)$:

$$\nabla_1 X_t = (1 - 0.7186_{0.0742}B)Z_t \text{ with } \hat{\sigma}_Z^2 = 0.2371$$

(1) Invertibility and Stationarity


For model (A):

$$(1 - 0.2373_{0.0963}B)\nabla_1 X_t = (1 - 0.8656_{0.0547}B)Z_t$$

Since $|\theta_1| < 1$ and $|\phi_1| < 1$, we can conclude that this model is both stationary and invertible. We can also plot the roots:

```
source('plot.roots.R')
plot.roots(polyroot(c(1, -0.2373)),polyroot(c(1, -0.8656)), main="Model A Roots ")
```

Model A Roots



From, the plot above, we can see that both roots are within the unit circle, which confirms our conclusion.


For model (B):

$$\nabla_1 X_t = (1 - 0.7186_{0.0742} B) Z_t$$

Since this is a pure MA model, it's automatically stationary. Also, $|\theta_1| < 1$, so we can conclude that this model is both stationary and invertible. We can also plot the roots:

```
source('plot.roots.R')
plot.roots(NULL, polyroot(c(1, -0.7186)), main="Model B Roots ")
```

Model B Roots





From, the plot above, we can see that the root is within the unit circle, which verifies our conclusion.


(2) Diagnostic Checking on Model (A): $ARIMA(1,1,1)$

```
op <- par(mfrow = c(2,2))
res.arima111 <- residuals(arima111)
hist(res.arima111, breaks = 20, xlab="", prob=TRUE, main = 'Histogram of ARIMA(1,1,1) Residuals')
m.arima111 <- mean(res.arima111)
std.arima111 <- sqrt(var(res.arima111))
curve(dnorm(x,m.arima111,std.arima111), add=TRUE )
plot.ts(res.arima111)
abline(h=mean(res.arima111), col="blue")
qqnorm(res.arima111,main= "Normal Q-Q Plot for Model A")
qqline(res.arima111,col="blue")
acf(res.arima111, lag.max=40)
```



Histogram of ARIMA(1,1,1) Residuals



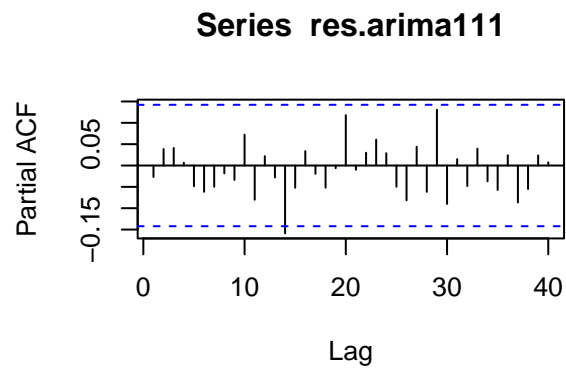
Normal Q-Q Plot for Model A



Series res.arima111



```
pacf(res.arima111, lag.max=40)  
par(op)
```



From the plots above, we can say that there is no trend, no visible change of variance, no seasonality, sample mean is almost zero, histogram and Q-Q plot all look good.

Now, we can apply tests to residuals:

```
shapiro.test(res.arima111)
```

```
##  
##  Shapiro-Wilk normality test  
##  
## data:  res.arima111  
## W = 0.99219, p-value = 0.4014
```

```
Box.test(res.arima111, type=c("Box-Pierce"), lag = 14, fitdf = 2)
```

```
##  
##  Box-Pierce test  
##  
## data:  res.arima111  
## X-squared = 9.3912, df = 12, p-value = 0.6692
```

```
Box.test(res.arima111, type=c("Ljung-Box"), lag = 14, fitdf = 2)
```

```
##  
## Box-Ljung test  
##  
## data: res.arima111  
## X-squared = 10.063, df = 12, p-value = 0.6104
```

```
Box.test((res.arima111)^2, type=c("Ljung-Box"), lag = 14, fitdf = 0)
```

```
##  
## Box-Ljung test  
##  
## data: (res.arima111)^2  
## X-squared = 15.916, df = 14, p-value = 0.3185
```

```
ar(res.arima111, aic = TRUE, order.max = NULL, method = c("yule-walker"))
```



```
##  
## Call:  
## ar(x = res.arima111, aic = TRUE, order.max = NULL, method = c("yule-walker"))  
##  
##  
## Order selected 0 sigma^2 estimated as 0.229
```

Since the residuals passed Shapiro test, Box-Pierce test, Ljung-Box test, Mcleod-Li test and AR order selects 0, it's safe to conclude that model A passed all the tests.


(3) Diagnostic Checking on Model (B): $MA(1)$

```
op <- par(mfrow = c(2,2))  
res.ma1 <- residuals(ma1)  
hist(res.ma1, breaks = 20, xlab="", prob=TRUE, main = 'Histogram of Model B Residuals')  
m.ma1 <- mean(res.ma1)  
std.ma1 <- sqrt(var(res.ma1))  
curve(dnorm(x,m.ma1,std.ma1), add=TRUE )  
plot.ts(res.ma1)  
abline(h=mean(res.ma1), col="blue")  
qqnorm(res.ma1,main= "Normal Q-Q Plot for Model B")  
qqline(res.ma1,col="blue")  
acf(res.ma1, lag.max=40)
```


Histogram of Model B Residuals




Normal Q-Q Plot for Model B



Series res.ma1



```
pacf(res.ma1, lag.max=40)
par(op)
```



From the plots above, we can say that there is no trend, no visible change of variance, no seasonality, sample mean is almost zero, histogram and Q-Q plot all look good.

Now, we can apply tests to residuals:

```
shapiro.test(res.ma1)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  res.ma1
## W = 0.99291, p-value = 0.4891
```

```
Box.test(res.ma1, type=c("Box-Pierce"), lag = 14, fitdf = 1)
```

```
##
##  Box-Pierce test
##
## data:  res.ma1
## X-squared = 14.06, df = 13, p-value = 0.3696
```

```
Box.test(res.ma1, type=c("Ljung-Box"), lag = 14, fitdf = 1)
```

```
##  
## Box-Ljung test  
##  
## data: res.ma1  
## X-squared = 14.9, df = 13, p-value = 0.3136
```

```
Box.test((res.ma1)^2, type=c("Ljung-Box"), lag = 14, fitdf = 0)
```

```
##  
## Box-Ljung test  
##  
## data: (res.ma1)^2  
## X-squared = 14.397, df = 14, p-value = 0.4206
```

```
ar(res.ma1, aic = TRUE, order.max = NULL, method = c("yule-walker"))
```

```
##  
## Call:  
## ar(x = res.ma1, aic = TRUE, order.max = NULL, method = c("yule-walker"))  
##  
##  
## Order selected 0 sigma^2 estimated as 0.2364
```

Since the residuals passed Shapiro test, Box-Pierce test, Ljung-Box test, Mcleod-Li test and AR order selects 0, it's safe to conclude that model B passed all the tests.

(4) Final Model Selection:

As we can see, both models passed all the test. According to the principle of parsimony, I should choose the one with the least coefficients, which is the $MA(1)$ model. Hence, the final model for de-trended data: X_t follows $MA(1)$ model: $\nabla_1 X_t = (1 - 0.7186_{0.0742} B) Z_t$ with $\hat{\sigma}_Z^2 = 0.2371$.

Forecasting

Now, we can forecast by using our model B:

```
library(forecast)  
# Fit into model and forecast  
fit <- arima(pek.ts, order=c(0,1,1), method="ML")  
forecast(fit)
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 2011	12.90245	12.27842	13.52649	11.94807	13.85683
## 2012	12.90245	12.25417	13.55073	11.91099	13.89391
## 2013	12.90245	12.23080	13.57410	11.87525	13.92965
## 2014	12.90245	12.20822	13.59668	11.84072	13.96419
## 2015	12.90245	12.18635	13.61855	11.80727	13.99764
## 2016	12.90245	12.16513	13.63978	11.77481	14.03009
## 2017	12.90245	12.14450	13.66041	11.74326	14.06164
## 2018	12.90245	12.12441	13.68049	11.71255	14.09236
## 2019	12.90245	12.10484	13.70006	11.68261	14.12230
## 2020	12.90245	12.08573	13.71917	11.65339	14.15152

To produce graph with 1 forecast on data:

```
pred<- predict(fit, n.ahead = 2)
```

```
U= pred$pred + 2*pred$se # Upper bound of prediction interval
```

```
L= pred$pred - 2*pred$se # Lower bound of prediction interval
```

Plot the forecast value for 2011 and 2012 with 95% CI on original data

```
ts.plot(pek.ts, xlim=c(1821,2012), ylim = c(min(pek.ts),max(U)), main = 'Forecast of Original Data Using
```

```
lines(2011:2012, y = U, col="blue", lty=2)
```

```
lines(2011:2012, y = L, col="blue", lty=2)
```

```
points(2011:2012, pred$pred, col="red")
```


```
legend("topleft",
```

```
  legend = c('Real Values', 'Forecasted Values', '95% CI'),
```

```
  fill = c('black','red','blue'),
```

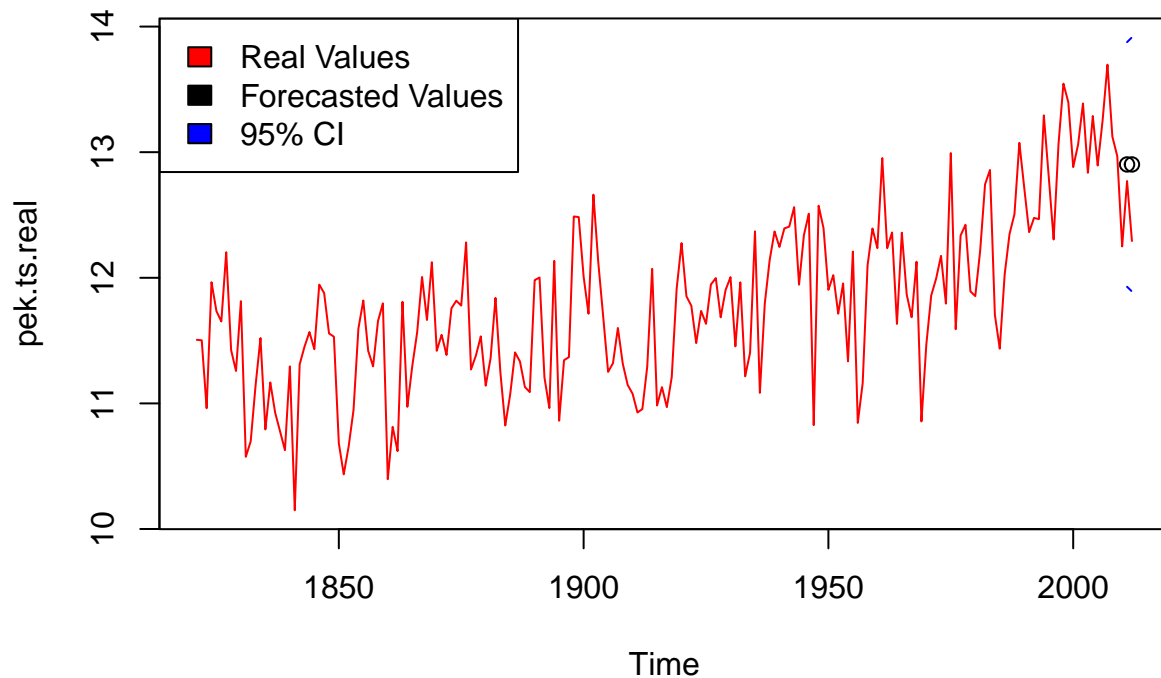
```
  border = "black")
```

Forecast of Original Data Using MA(1) Model



```
# Plot the forecast value for 2011 and 2012 with 95% CI on original data with real values
pek.ts.real <- ts(data[,2], start = 1821, frequency = 1)
ts.plot(pek.ts.real, xlim = c(1821,2012), ylim = c(min(pek.ts),max(U)), col="red", main = 'Forecast of Original Data Using MA(1) Model')
lines(U, col="blue", lty="dashed")
lines(L, col="blue", lty="dashed")
points(2011:2012, pred$pred, col="black")
legend("topleft",
      legend = c('Real Values', 'Forecasted Values', '95% CI'),
      fill = c('red','black','blue'),
      border = "black")
```


Forecast of Original Data Using MA(1) Model with Real Values



Look at the plots above, we can see that the real values from our test set is within the prediction interval. And we can plot the zoom-in graphs to see it more clearly:

```
# Plot the forecast value for 2011 and 2012 with 95% CI on original data: Zoom In
ts.plot(pek.ts, xlim=c(1950,2012), ylim = c(min(pek.ts),max(U)), main = 'Forecast of Original Data Using MA(1) Model')
lines(2011:2012, y = U, col="blue", lty=2)
lines(2011:2012, y = L, col="blue", lty=2)
points(2011:2012, pred$pred, col="red")
legend("topleft",
      legend = c('Real Values', 'Forecasted Values', '95% CI'),
      fill = c('black','red','blue'),
      border = "black")
```


Forecast of Original Data Using MA(1) Model



```
# Plot the forecast value for 2011 and 2012 with 95% CI on original data with real values: Zoom In
ts.plot(pek.ts.real, xlim = c(1950,2012), ylim = c(min(pek.ts),max(U)), col="red", main = 'Forecast of 
lines(U, col="blue", lty="dashed")
lines(L, col="blue", lty="dashed")
points(2011:2012, pred$pred, col="black")
legend("topleft",
      legend = c('Real Values', 'Forecasted Values', '95% CI'),
      fill = c('red','black','blue'),
      border = "black")
```

Forecast of Original Data Using MA(1) Model with Real Values

